

Shreyash Kondakindi

(858) 370-3500 | kshreyashreddy@gmail.com | [LinkedIn](#) | shreyashreddyk.github.io | San Diego, CA

Summary

Applied AI Engineer building GenAI systems and distributed data platforms, from model evaluation and serving to Spark/Dataproc pipelines and user-facing AI applications. Experience spans healthcare, finance, and university operations with end-to-end delivery across FastAPI/React products, cloud orchestration, analytics infrastructure, and production-minded observability.

Education

University of California, San Diego <i>M.S. Data Science, GPA: 3.94/4.0</i>	San Diego, CA <i>2024 – 2026</i>
Birla Institute of Technology and Science, Pilani <i>B.E. Mechanical Engineering, Minor in Finance, GPA: 9.2/10</i>	Hyderabad, India <i>2019 – 2023</i>

Work Experience

University of California, San Diego <i>Graduate Student Researcher</i>	San Diego, CA <i>Sept 2025 – Jan 2026</i>
--	---

- Researched SLO-aware LLM serving and scheduling with Prof. Umesh Bellur; co-authored an ACM Computing Surveys manuscript on batching, routing, and latency-throughput trade-offs.
- Complemented the research focus with structured extraction post-training using SFT/QLoRA, DPO, Pydantic validation, and evaluation harnesses, improving Micro F1 from 0.303 to 0.773 while maintaining 0.998 schema pass.
- Benchmarked vLLM serving on NVIDIA L4 across mixed prompt workloads, identifying a throughput knee at concurrency 16–32 and p99 expansion from 6.2s to 21.7s while eliminating token-limit failures with dynamic context budgeting.

University of California, San Diego <i>GenAI Analyst – BFS BI & Financial Reporting</i>	San Diego, CA <i>Jun 2025 – Sept 2025</i>
---	---

- Automated AI-assisted remediation of 500+ broken links and self-serve analytics workflows across BFS BI and Financial Reporting, reducing manual maintenance 30%.
- Analyzed 40K+ service tickets with Python, SQL, and Oracle AI to prioritize issue patterns, improve triage quality, and reduce MTTR 20%.
- Built OTBI and Cognos dashboards that surfaced queue health and operational bottlenecks for proactive issue resolution.

University of California, San Diego <i>Big Data Analytics Engineer</i>	San Diego, CA <i>Mar 2025 – Jun 2025</i>
--	--

- Ran Spark and Spark SQL pipelines on GCP Dataproc over 12M+ Parquet weather records, using custom UDFs and RDD transforms to convert packed annual measurements into station-day climate series for distributed analysis.
- Scheduled Dataproc jobs with Cloud Composer/Airflow, persisted Parquet checkpoints for fault-tolerant ETL, and engineered a 9-feature PCA representation from 366 daily signals for downstream climate modeling.
- Trained XGBoost to 72.8% test accuracy on Köppen climate classification.

Accordion <i>Data Scientist – Healthcare Analytics</i>	Virginia Beach, VA Hyderabad, India <i>July 2023 – June 2024</i>
--	--

- Orchestrated a 2-month cloud migration of 200GB+ HIPAA-compliant data, leading a 4-person analyst/engineering transition and scheduling Snowflake Medallion pipelines across 3 source systems for RCM, finance, and operations reporting.
- Automated Alteryx, Python, API, and Tableau workflows, saving 8+ client hours daily, cutting manual processing by 90 minutes per day, and improving query performance 30%+.
- Built CV-assisted healthcare review with YOLOv8 localization, SAM 2.1 segmentation, and Grounding DINO/SAM annotation support, reducing manual review 80% and labeling from 5 minutes to under 1 minute.

J.P. Morgan & Co. <i>Data Science Intern – Corporate & Investment Bank Analytics</i>	Mumbai, India <i>Jan 2023 – June 2023</i>
--	---

- Engineered Python pipelines over AlphaSense, Eikon, Bloomberg, SEC XMLs, and broker PDFs into Azure SQL for banker-facing M&A and risk research across 15+ sectors.
- Refactored DCF, comps, cash-flow, and ownership workflows with pandas, NumPy, Excel, and VBA, reducing manual reconciliation 40% and supporting a 30-page Power BI-backed sell-side pitch for a U.S. grocery retail client.

Projects

<u>Prosody – Fullstack AI Coaching Workspace</u>	Apr 2026 – Jun 2026
---	----------------------------

- Built and deployed a 2-service AI coaching workspace with React 18, FastAPI, Supabase Auth/Postgres/Storage, and Cloud Run for authenticated uploads, persistent sessions, and 8+ user-scoped artifacts including summaries, flashcards, sources, and coaching metrics.
- Integrated voice-agent workflows with Pipecat, Deepgram Flux STT, OpenAI, and ElevenLabs for realtime coaching flows during development and diagnostics.
- Operationalized delivery with GitHub Actions CI, Workload Identity Federation, Secret Manager, JWT-scoped persistence, and health/readiness/meta controls across both services.

<u>SEC Filing Intelligence Copilot – RAG & Evaluation</u>	Mar 2026 – May 2026
--	----------------------------

- Indexed 10 live 10-K and 10-Q filings into a 1,611-chunk corpus across 5 semiconductor companies using FastAPI, Chroma, hybrid retrieval, and reranking for citation-grounded analyst workflows.
- Added deterministic evaluation and observability with Docker, GKE, and Prometheus/Grafana; validated hit_rate@4=1.0, recall@4=1.0, MRR=0.9167, citation validity=1.0, and abstention accuracy=1.0.

Skills
AI / ML: Python, SQL, PyTorch, XGBoost, Hugging Face, OpenAI API, Deepgram, ElevenLabs, vLLM, Pydantic
Platforms / Data: FastAPI, React, Docker, Cloud Run, GKE, Supabase, Spark, Dataproc, Cloud Composer/Airflow, Snowflake, Azure SQL, Chroma, Matillion, Alteryx, Tableau, Power BI, GitHub Actions, Prometheus/Grafana, Terraform