

Shreyash Kondakindi

+1 (858) 370-3500 | shreyashkondakindi16@gmail.com | [LinkedIn](#) | [Portfolio](#) | San Francisco, CA (Open to relocation)

Summary

Applied AI Engineer with 2+ years of experience and building GenAI systems, distributed data platforms, and user-facing AI applications. Experience spans model evaluation and serving, Spark/Dataproc pipelines, FastAPI/React delivery, cloud orchestration, analytics infrastructure, and production-minded observability across healthcare, finance, and university operations.

Skills

Programming & ML: Python, SQL, PyTorch, scikit-learn, XGBoost, Hugging Face, pandas, NumPy, Pydantic

GenAI & Agents: RAG, OpenAI API, LangChain, LangGraph, Google ADK, Chroma, Ragas, vLLM, Pipecat

Fullstack & Data: TypeScript, React, FastAPI, Supabase/PostgreSQL, Spark, Dataproc, Composer/Airflow, Snowflake

Cloud & MLOps: GCP (Cloud Run, GKE), Docker, Git/GitHub Actions, Terraform, Prometheus/Grafana, pytest, CI/CD

Work Experience

Graduate Student Researcher

University of California, San Diego

San Diego, CA

Sept 2025 – Jan 2026

- Researched SLO-aware LLM serving and scheduling with Prof. Umesh Bellur; co-authored an ACM Computing Surveys [manuscript](#) on batching, routing, and latency-throughput trade-offs.
- Built structured extraction post-training and evaluation workflows with SFT/QLoRA, DPO, Pydantic validation, and benchmark harnesses, improving Micro F1 from 0.303 to 0.773 while maintaining 0.998 schema pass.
- Benchmarked vLLM serving on NVIDIA L4 across mixed prompt workloads, identifying a throughput knee at concurrency 16–32 and p99 expansion from 6.2s to 21.7s while eliminating token-limit failures with dynamic context budgeting.

GenAI Analyst – BFS BI & Financial Reporting

University of California, San Diego

San Diego, CA

Jun 2025 – Sept 2025

- Automated remediation of 500+ broken links across BFS BI and Financial Reporting, cutting manual maintenance 30% and improving reliability for self-serve analytics.
- Analyzed 40K+ service tickets with Python and Oracle AI, cutting triage MTTR 20% while prioritizing recurring issue patterns.
- Built OTBI and Cognos dashboards that surfaced queue health and operational bottlenecks for proactive issue resolution.

Big Data Analytics Engineer

University of California, San Diego

San Diego, CA

Mar 2025 – Jun 2025

- Ran Spark and Spark SQL pipelines on GCP Dataproc over 12M+ Parquet weather records, using custom UDFs and RDD transforms to convert packed annual measurements into station-day climate series.
- Scheduled Dataproc jobs with Cloud Composer/Airflow, persisted Parquet checkpoints for fault-tolerant ETL, and engineered a 9-feature PCA representation from 366 daily signals for downstream climate modeling.
- Trained XGBoost to 72.8% test accuracy on Köppen climate classification.

Data Scientist – Healthcare Analytics

Accordion

Virginia Beach, VA | Hyderabad, India

July 2023 – June 2024

- Led a 2-month cloud migration of 200GB+ HIPAA-compliant data, guiding a 4-person analyst/engineering transition and scheduling Snowflake Medallion pipelines across 3 source systems for RCM, finance, and operations reporting.
- Automated Alteryx, Python, API, and Tableau workflows, saving 8+ client hours daily, cutting manual processing by 90 minutes per day, and improving query performance 30%+.
- Built CV-assisted healthcare review with YOLOv8 localization, SAM 2.1 segmentation, and Grounding DINO/SAM annotation support, reducing manual review 80% and labeling from 5 minutes to under 1 minute.

Data Science Intern – Corporate & Investment Bank Analytics

J.P. Morgan & Co.

Mumbai, India

Jan 2023 – June 2023

- Designed Python pipelines over AlphaSense, Eikon, Bloomberg, SEC XMLs, and broker PDFs into Azure SQL to support banker-facing M&A and risk research across 15+ sectors.
- Created DCF, comps, cash-flow, and ownership workflows with pandas, NumPy, Excel, and VBA, reducing manual reconciliation 40% and supporting a 30-page Power BI-backed sell-side pitch for a U.S. grocery retail client.

Projects

Prosody – Fullstack AI Coaching Workspace | [GitHub](#)

Apr 2026 – Jun 2026

- Built a 2-service AI coaching workspace with React 18, FastAPI, Supabase Auth/Postgres/Storage, and Cloud Run for authenticated uploads, persistent sessions, and 8+ user-scoped artifacts including summaries, flashcards, sources, and coaching metrics.
- Added voice-agent workflows with Pipecat, Deepgram Flux STT, OpenAI, and ElevenLabs for realtime coaching turns, streaming transcription, and spoken feedback during development.
- Deployed both services with GitHub Actions CI, Workload Identity Federation, Secret Manager, JWT-scoped persistence, and health/meta endpoints for release checks.

SEC Filing Intelligence Copilot – RAG & Evaluation | [GitHub](#)

Mar 2026 – May 2026

- Indexed 10 live 10-K and 10-Q filings into a 1,611-chunk corpus across 5 semiconductor companies using FastAPI, Chroma, hybrid retrieval, and reranking for citation-grounded analyst workflows.
- Validated evals with Docker, GKE, and Prometheus/Grafana, achieving recall@4=1.0, hit_rate@4=1.0, MRR=0.9167, and perfect citation validity and abstention.

Education

Masters in Data Science, GPA: 3.94/4.0

University of California, San Diego

San Diego, CA

2024 – 2026

Bachelors in Engineering, Minor in Finance, GPA: 3.75/4.0

Birla Institute of Technology and Science, Pilani

Hyderabad, India

2019 – 2023